

Manuscript JSS2717
“sparseHessianFD: An R Package for Estimating Sparse Hessian Matrices”

Response to reviewers

I would like to express my sincere appreciation for the time and effort that the reviewer invested in the review. It is obvious to me that the reviewer went through the paper and code in great detail. In particular, I want to thank the reviewer for bringing the complex step method to my attention. I *never* use complex numbers or functions in my own work, so I typically ignore anything with the word "complex" in it. But this review prompted me to learn something new. That is always fun.

I am happy to write that I was able to incorporate a large majority of the reviewer's suggestions in the revised paper and package. For the others, I hope that the editor and reviewer will accept my explanations and excuses. Of course, if there are any strong objections to how I responded to the reviewer's concerns, I am happy to make further changes in advance of publication.

0. *I had a minor problem checking the package tarball:*

(error messages related to algorithm.sty removed).

I don't remember if there is an easy way to distribute "special" sty files, or if you must force the user to install them. If the sty is not really critical it might be better to omit it.

I suspect that the error is a result of the *algorithm* package not being present in the reviewer's LaTeX installation. I would like to continue to use this package, because it does allow for nicely formatted algorithms. I believe that this is a relatively common package, and one that should be easily installed through CTAN. However, if there is a strong objection, I may be able to find another way to format the algorithms.

1. *I am not familiar with reference classes, but for most users I think it should not be important that the package uses them. However, as a user, I find it disconcerting that sparseHessianFD() is used as a (constructor) function (eg. p12, and in the example in ?sparseHessianFD) but is not documented in the help with 'usage' and 'arguments' as is usual for functions. There seems to be little guidance in the usual R places about how reference classes should be documented. Perhaps you could seek some guidance from the community on this point.*

Note to me: issue here is that there is no documentation for `sparseHessianFD()` as a function. Good idea to create a documentation page for that.

2. *It is pointed out (p2) that requirements are burdensome and emphasized that the package is not appropriate for all uses. Also, the conditions 1-5 are fairly stringent. The reader is left with the impression that there may be few applications. I think the value of the package could be promoted better. Perhaps a small list of different types or applications of hierarchical problems could be given, or an indication of other classes of problems where the conditions would apply.*

I included this section primarily as a way to avoid overselling the method, but the point that the reader might be scared away by this list of restrictions is well-taken. In the revision, I removed the list of conditions, and wove them into the text. The result takes a more positive tone than in the original submission.

3. *Condition 5 might be relaxed to a local condition in a neighborhood of the evaluation point, rather than a global condition, but I do not know if that would have any practical value.*

The list of conditions is no longer in the paper, but to this point, I added a definition of a “structural zero” in the third paragraph. The sparsity pattern represents the structure of the Hessian, not its value, so any element that *could* be zero, but does not have to be zero, is not a structural zero. Thus, the condition is global and not local.

Note that other than a potential loss of efficiency, there is no disadvantage to erring on the side of an element being non-zero.

4. *Section 2.1 p7 I think could be made more easily readable by adding a few more hints about dimension, and more care about use of the term "coefficients vector", for example:*
- *'continuous covariates x_i ,' -> 'continuous covariates $x_i \in \mathbb{R}^k$,*
 - *'heterogeneous coefficient vector β_i ' -> 'heterogeneous coefficient vector $\beta_i \in \mathbb{R}^k$,*
 - *'The coefficients are distributed' -> 'The coefficient vectors β_i are distributed'*

Thank you for the suggestion. Changes have been made throughout the paper.

5. *p7 following "the cross-partial derivatives" $H_{\beta_i, \beta_j} = D_{\beta_i, \beta_j}^2 = 0$ for all $i \neq j$.*

Is this a definition of H_{β_i, β_j} in terms of D^2 or a statement of equality?

Also, I am confused by the single subscript β_i to

hess here, but a double subscript just above in "Thus, $H_{\beta_{ik}, \mu_k} \neq 0$ ".

Based on this comment, I rewrote Section 2.2 to distinguish the cross-partial derivatives D^2 from the corresponding elements of the Hessian H . I did this by defining x_i as a subset of the elements of x , and $\text{Ind}(x_i)$ as an “indexing” function. The revised version is slightly more formal and precise, and I hope that this is the kind of change you had in mind. That said, I do not think it makes the paper easier to read, and, with your permission, I would prefer to revert to the simpler language in the original manuscript.

6. *I think possibly (16) is "banded" and (17) "block arrow" rather than the reverse which is indicated.*

No, the descriptions of the two sparsity patterns are correct. The first index in the subscript of β refers to a household, and the second index refers to a covariate. Thus, $D_{\beta_{11}, \beta_{12}}^2 \neq 0$, but $D_{\beta_{11}, \beta_{21}}^2 = 0$. When the second index changes before the first, the pattern is “block-arrow.” When the first index changes before the second, the pattern is “banded.”

7. *Code in the paper prior to table 4 needs to be copied from the vinettes/sparseHessianFD.Rnw file. For those trying to reproduce results it would be nice if this were mentioned in the file replication.R.*

8. *p9, last line. The $R>$ at the beginning of*

$R> \text{obj} <- \text{sparseHessianFD}(x, \text{fn}, \text{gr}, \text{rows}, \text{cols}, \dots)$

suggests that this is code that can be entire at the command line, but ... causes an error when entered (and the arguments have not been defined at this point in the vignette. Instead it should be indicated as the usage syntax.

I removed the R> prompt, and now refer to that code as a usage syntax.

9. *p10. "where ... represents all other named arguments" I think the usual usage is the ... represents the arguments other than the "named" ones, so it is probably better to just say "other arguments".*

Done.

10. *On p12, the function calls `all.equal(f, true.f)` and `all.equal(gr, true.grad)` are comparing `f` and `gr` calculated with the exact calculation, so the difference is zero:*

Code removed

It is not clear to me whether `obj$fn()` and `obj$gr()` use code as in the true functions or a modified version using sparse techniques. Some further clarification would be helpful.

On the other hand, `all.equal(hs, true.hess)` is comparing a true analytic calculation with a first order simple difference approximation using the true gradient function:

```
max(abs( obj$hessian(P)
        - binary.hess(P, data=binary, priors=priors, order.row=order.row)))
[1] 2.786891e-06
```

which might also be mentioned in the text. (Really just for exposition purposes, after all, it is almost the main purpose of the package.)

Good point. I revised that section to explain that evaluations of the function and gradient have to be identical to the true values, and I even replaced `all.equal` with `identical` to highlight that fact. I then explained that the `sparseHessianFD` calculation of the Hessian will not be identical to the true value, and report the mean relative difference. I chose the mean relative difference over maximum absolute difference because the former is what the `all.equal` function uses to test if two matrices are equal within numeric tolerance.

11. *I think it would be instructive to add some of the following comparisons with the above on p12. The package `numDeriv` function `hessian` by default does a second order Richardson approximation using the true function value approximation. This involves a very large number of function evaluations in an attempt to obtain some accuracy, but the accuracy is limited by being an approximation of a second difference:*

Hessian Richardson

Since the hessian is the first difference of the gradient, which is the calculation used by `obj$hessian()` in `sparseHessianFD`, one could also use the function `numDeriv::jacobian`:

Jacobian Richardson

This is still doing the calculation intensive Richardson approximation. The calculation which would seem to most closely resemble what is done by `obj$hessian()` is

`Jacobian simple`

Another very interesting comparison is

`Jacobian complex`

The complex step derivative provides extremely accurate approximations with a number of function evaluation similar to the simple method. (This does not seem to be anticipated by footnote 1 in the paper.) However, the method imposes some serious requirements on the function. (Something like complex analytic even though the user may only be interested in the real part.) The code also has to accept complex arguments and return the complex result. Fortunately most R primitive work with complex numbers so the code requirement may happen accidentally, which can be partly verified by

`binary.grad(P + 0+1i, data=binary, priors=priors, order.row=order.row)`

returning a complex result. (This does not rule out all possible problems.)

As I recall, sums, multiplication, and exponentiation are all complex analytic, so it would not be too surprising if the example in the paper is too, but I have not analyzed that. However, based on the result being very good, it seems highly likely.

See response to Item 12

12. *A possible extension to the package would be to implement the complex method in the sparse code. The function `numDeriv::jacobian.default` implements both simple and complex, so provides a good comparison of the necessary (non-sparse) computation.*

Here I am combining my responses to items 11 and 12 together.

The complex step function is now implemented in the package, and is introduced in Section 2.5 of the paper. You are correct that “most R primitives work with complex numbers,” but there are many commonly used functions that do not, including `log1p`, `gamma`, and the probability distribution functions. This limitation on the kinds of functions that can be implemented easily in R, combined with the requirement that the function be holomorphic, and my own general unfamiliarity with the method, make me hesitant to promote it too much in the paper. But it is a nifty trick, and including it in the package was a good idea (thanks!).

In terms of accuracy, the paper includes comparisons between the true Hessian and the two `sparseHessianFD` methods (finite differences and complex step). I am not comparing accuracy with `numDeriv`, because that really is not the point of this paper. However, for the timing comparisons in Table 4, I did change the baseline `numDeriv` method from `hessian` with the Richardson method, to `jacobian` with the `complex` method. Removing the second-order approximation saves a lot of time, and lets me compute Table 4 without resorting to parallel computation. That will simplify attempts at replication, and removes another variable that might influence results.

13. *p.13 l. -7 Figure 3b should be Figure 3c*

Fixed. Thanks.

14. *The file replication.R does not set the RNG seed. This may not be too important if only times are generated, but will be if resulting values are included.*

I added a `set.seed` statement near the top of *replication.R*.

15. *Table 4 and 5. Some OS, processor, and memory details are helpful to put timing results in context.*

16. *Table 4 is really not the proper comparison. I think a comparison with `numDeriv::jacobian(binary.grad, P, method="simple", data=binary, priors=priors, order.row=order.row)`*

really serves to highlight the improvement of the sparse calculation because it is a valid comparison. Even though the results are not as exaggerated, they are still important:

(My laptop is a Intel(R) Core(TM) i5-3337U CPU @ 1.80GHz, 4GB RAM, SSD swap, running Mint variant of Ubuntu 14.04.2 LTS.)

Code and output removed

I agree. See my response to Item 12. I include the times from the `complex` method instead of `simple` because of the improved accuracy. I considered including both times in the table, but that would have cluttered the table, there was very little difference in the times between them, and the comparison between two `numDeriv` methods is not the point of this paper. But I completely agree that either `jacobian` method is a more valid comparison than using `hessian`.

17. *It is possible to do a larger example with this comparison:*

(On my laptop the next took about 30 hrs of which 24 was for the last, $N=2500$, $k=8$ comparison.)

```
cases_tab4b5 <- expand.grid(k=c(2,4,8),
                           N=c(10, 100, 1000, 2500))
runs_tab4b5 <- plyr::adply(cases_tab4b5, 1, run_test_tab4b, reps=20,
                          .parallel=run.par)

tab4b5 <- mutate(runs_tab4b5, ms=bench.time/1000000) %>%
  select(-bench.time) %>%
  spread(bench.expr, ms) %>%
  gather(method, hessian, c(numDeriv, sparse)) %>%
  mutate(M=N*k+k, hessian.df=hessian/df) %>%
  gather(stat, time, c(hessian, hessian.df)) %>%
  group_by(N, k, method, M, stat) %>%
  summarize(mean=mean(time), sd=sd(time)) %>%
  gather(stat2, value, mean:sd) %>%
```

```
dcast(N+k+M~stat+method+stat2,value.var="value") %>%
  arrange(M)
```

Output removed

I think the complex step takes a similar amount of time, but produces a more accurate result.

18. p.16 l. -7 "to to compute" -> "to compute"

Fixed. Thanks.

19. *While that package is useful, and reasonably demonstrated in the paper, I think it would be nice to expand the paper in some ways that might be deemed more "original research". Some possibilities are:*
- *Explain and try to assess how much of the speedup is due to simple sparseness and how much is due to the "sparse pattern" (p2) allowing for perturbing multiple elements together. (I think these are related but slightly different?)*
 - *Try to assess how much of the speedup is due to reduce computing demand and how much is due to different memory demand. (I had the impression in the larger problems with numDeriv that my computer started to use swap space, which resulted in a big slow down.)*
 - *Consider implementing a complex step method, and do a comparison.*
 - *Assess the difference when multiple CPUs are used. (run.par == FALSE vs run.par == TRUE)*

These are all good suggestions. Let me address them point by point.

- It really is all about the sparsity pattern, and specifically, the how few partitions are needed to group the variables. Of course, it will be hard to keep the number of partitions small if there are too many non-zero elements, so the sparsity itself has some effect. The arrangement of the non-zeros is related to the ordering of the variables, but the permutation handles that during the initialization of the *sparseHessianFD* argument. For example, variables arranged in a way that the Hessian has a block-arrow pattern are permuted to get a Hessian with a banded pattern instead. This is all discussed in the "Partitioning the variables" section of the paper.
- I have no idea how to do this in a generalizable way. I think it has a lot to do with the architecture of the computer itself. Figure 4 does break down the components of the computation into the time it takes to reserve memory (initialization), partition the variables, and compute gradients themselves.
I do not want to ignore this suggestion, but I am simply not clear on what is being asked. I am open to more specific suggestions.
- The complex step method is now supported, and comparisons of accuracy and computation time are now included in the paper.
- Parallel computing on a shared memory architecture really muddles up timing comparisons. The *only* reason I generated Table 4 in parallel was because of the *numDeriv hessian*

function. Now that `hessian` was replaced with `jacobian`, the baseline measures from *numDeriv* run fast enough that I can generate Table 4 in serial.

Your point is well taken that timing comparisons can be affected by whether computation is being run on multiple processors, but the “serialization” of Table 4 makes the issue moot.